

oueees-201906 talks Part 3/3: Centralized and distributed systems

Kenji Rikitake

25-JUN-2019

School of Engineering Science
Osaka University
Toyonaka, Osaka, Japan

@jj1bdx

Copyright ©2018-2019 Kenji Rikitake.
This work is licensed under a [Creative Commons](#)
[Attribution 4.0 International License](#).

Lecture notes

- <https://github.com/jj1wdx/oueees-201906-public/>
- Check out the README.md file and the issues!

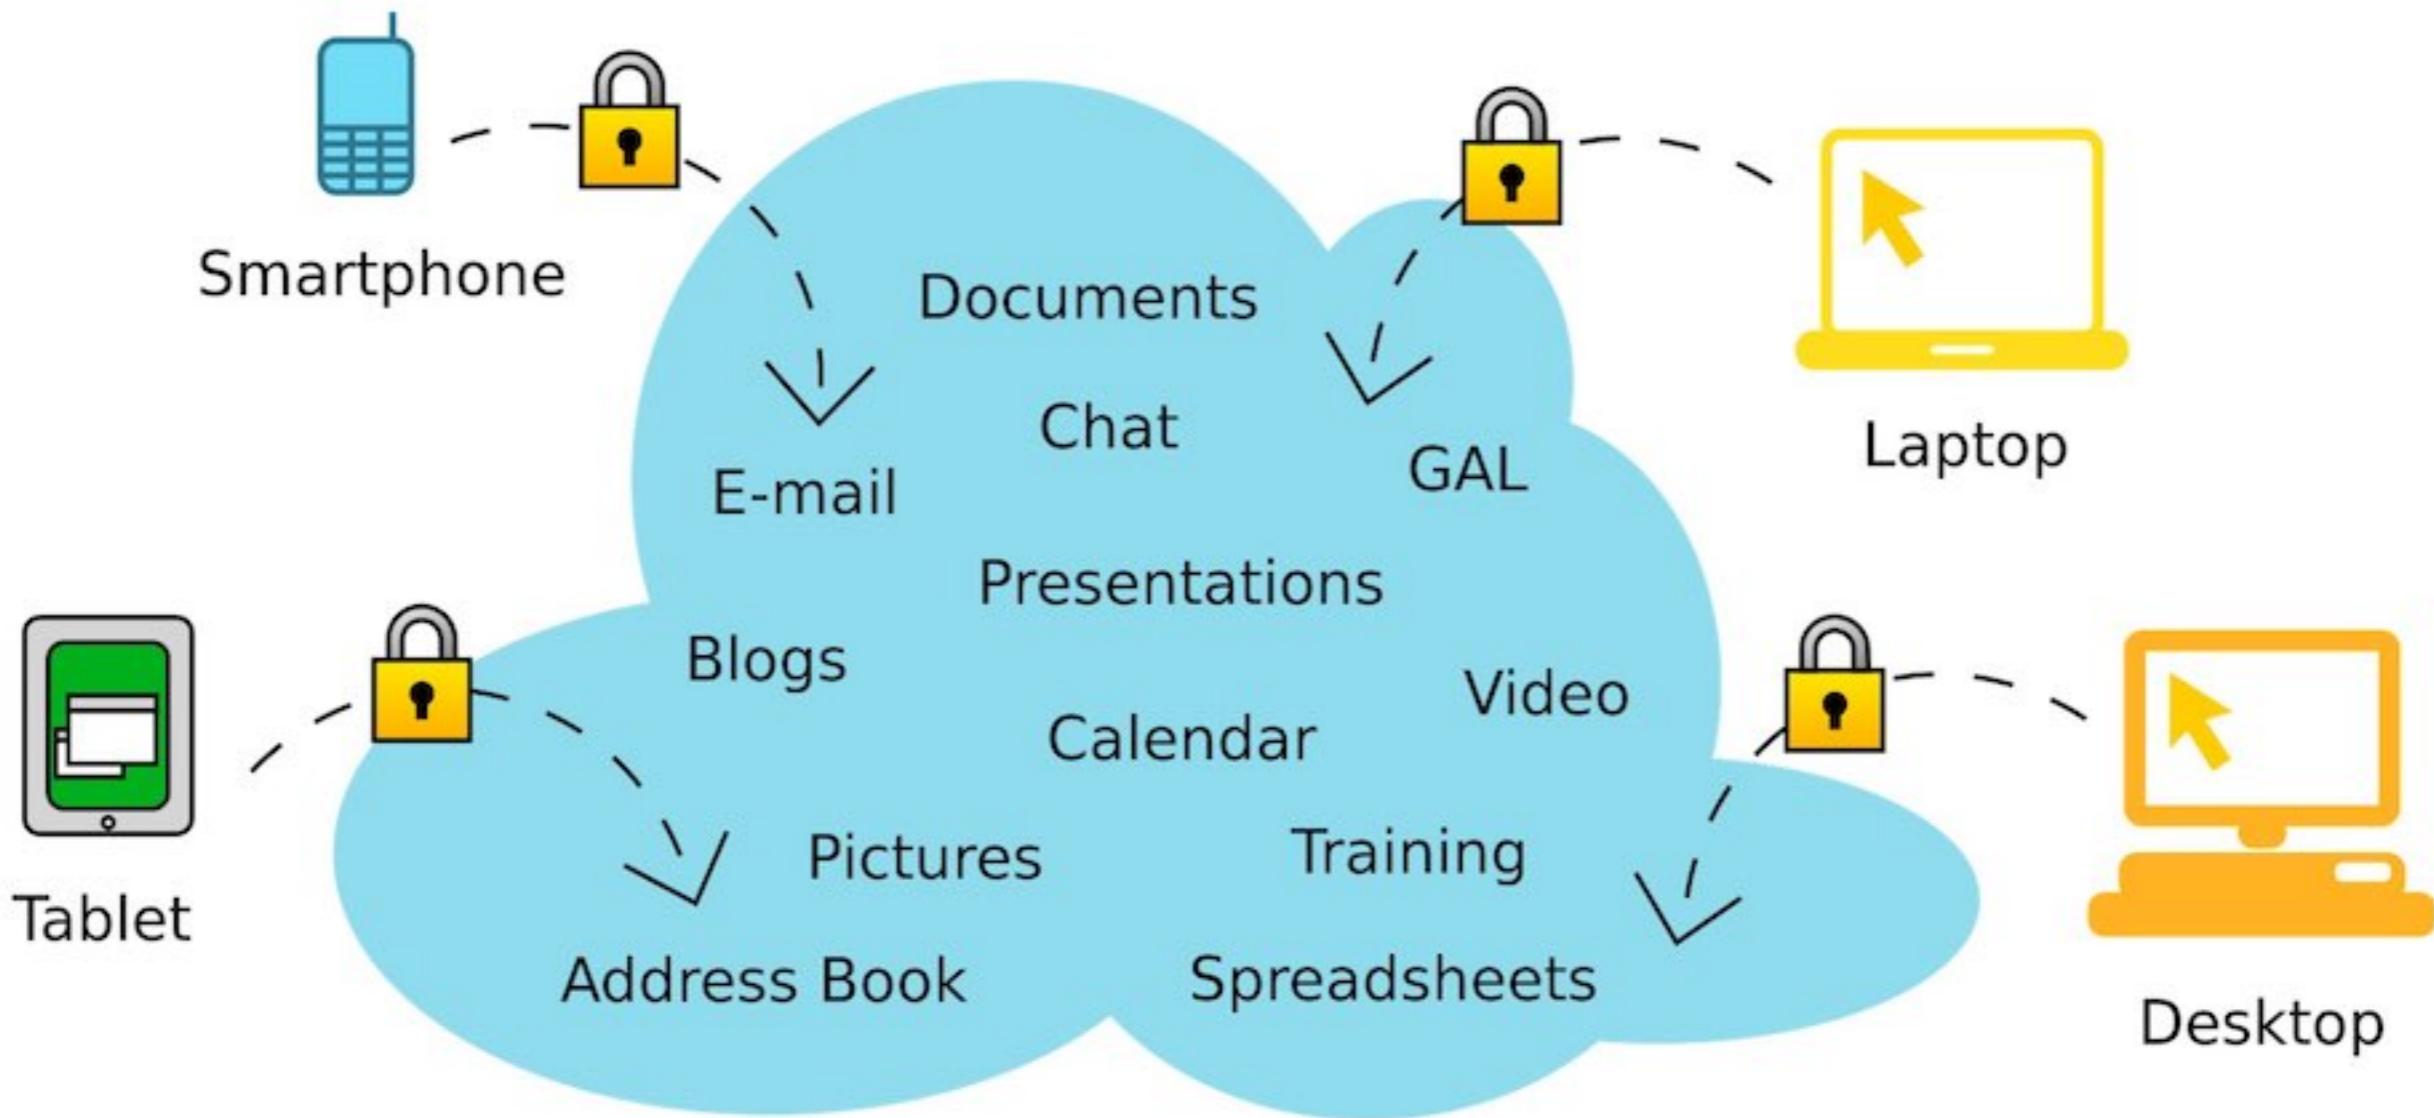
Reporting

- Keyword at the end of the talk
- URL for submitting the report at the end of the talk

Today's topic: centralized and distributed systems

A black and white photograph of a computer keyboard, a mug, and a smartphone on a desk. The keyboard is in the foreground, angled towards the bottom left. A dark mug is positioned behind it, and a smartphone with a cable is resting on the keyboard. In the background, a portion of a computer monitor is visible.

Modern computing *is* cloud computing



Cloud Computing

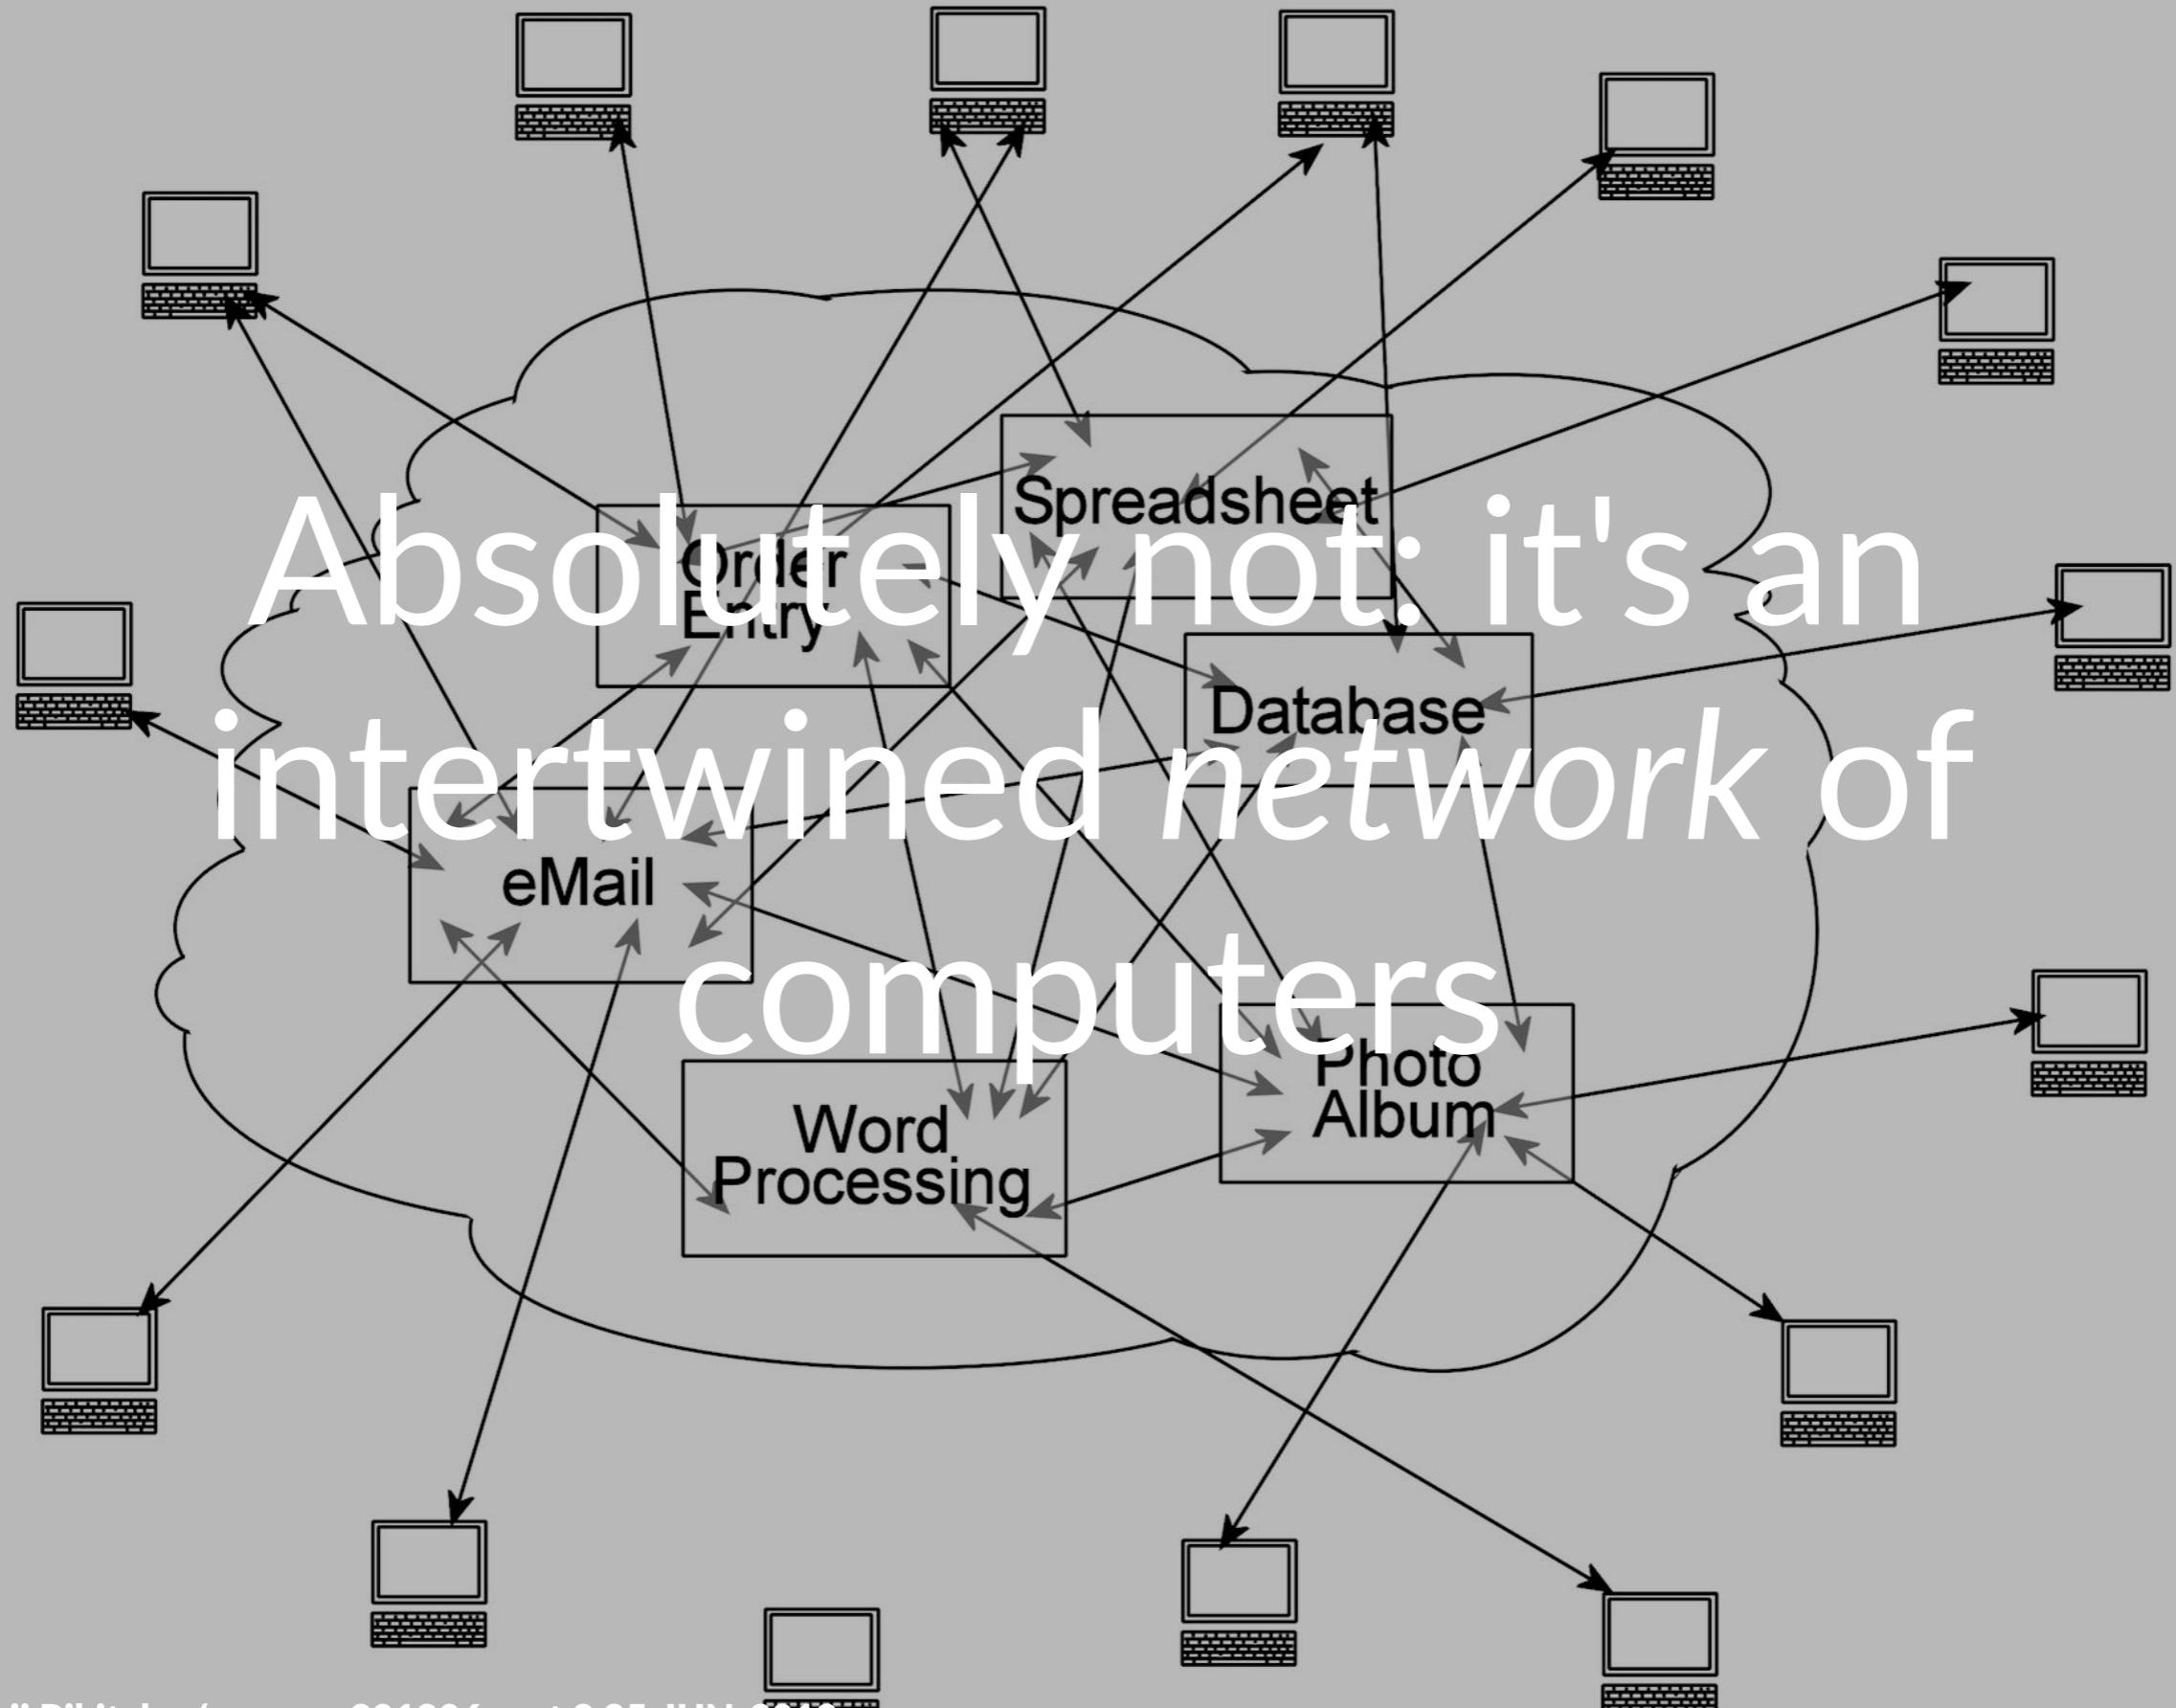
Having secure access to all your applications and data from any network device

Is cloud really a uniform and single entity?



Cloud Computing

Having secure access to all your applications and data from any network device



Web services are clusters of computers and networks

Thousands or millions of servers connected together


A physical server is separated into multiple virtual machines



Cloud computing is
centralization

An example of cloud computing: Internet of Things (IoT) and telemetering

“Things” or devices



Collecting data

Telemetering

- Mostly unidirectional (not really *the true and genuine internet*)
- Sensors/devices gathering data through internet and feed them to the servers in the cloud computing platforms
- *The servers compute*
- *Extremely centralized*

The social implication of cloud computing


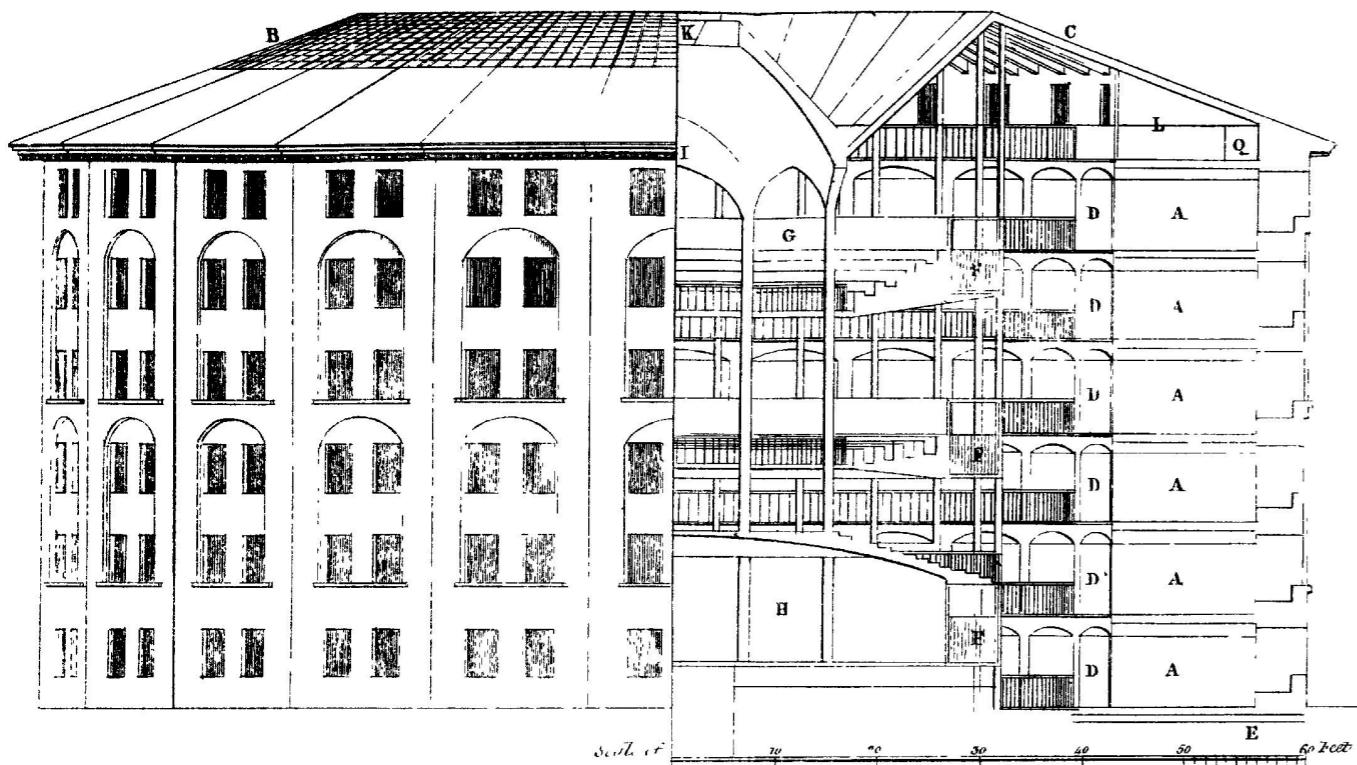
Centralized social behavior accelerated by cloud computing

- Sharing *everything* - no privacy
- *Panopticon*¹ style of governance, filtering, censorship, or autocracy
- Complete *externalization* of resources, leading to no personal control



¹ n. a circular prison with cells arranged around a central well, from which prisoners could at all times be observed. (New Oxford American Dictionary, Apple macOS 10.13.6)

Presidio Modelo: a panopticon prison





CONTRASTED RESIDENCES FOR THE POOR



INGSOC: the slogans ²

- War is peace
- Freedom is slavery
- Ignorance is strength
- Independent thinking = *thoughtcrime*

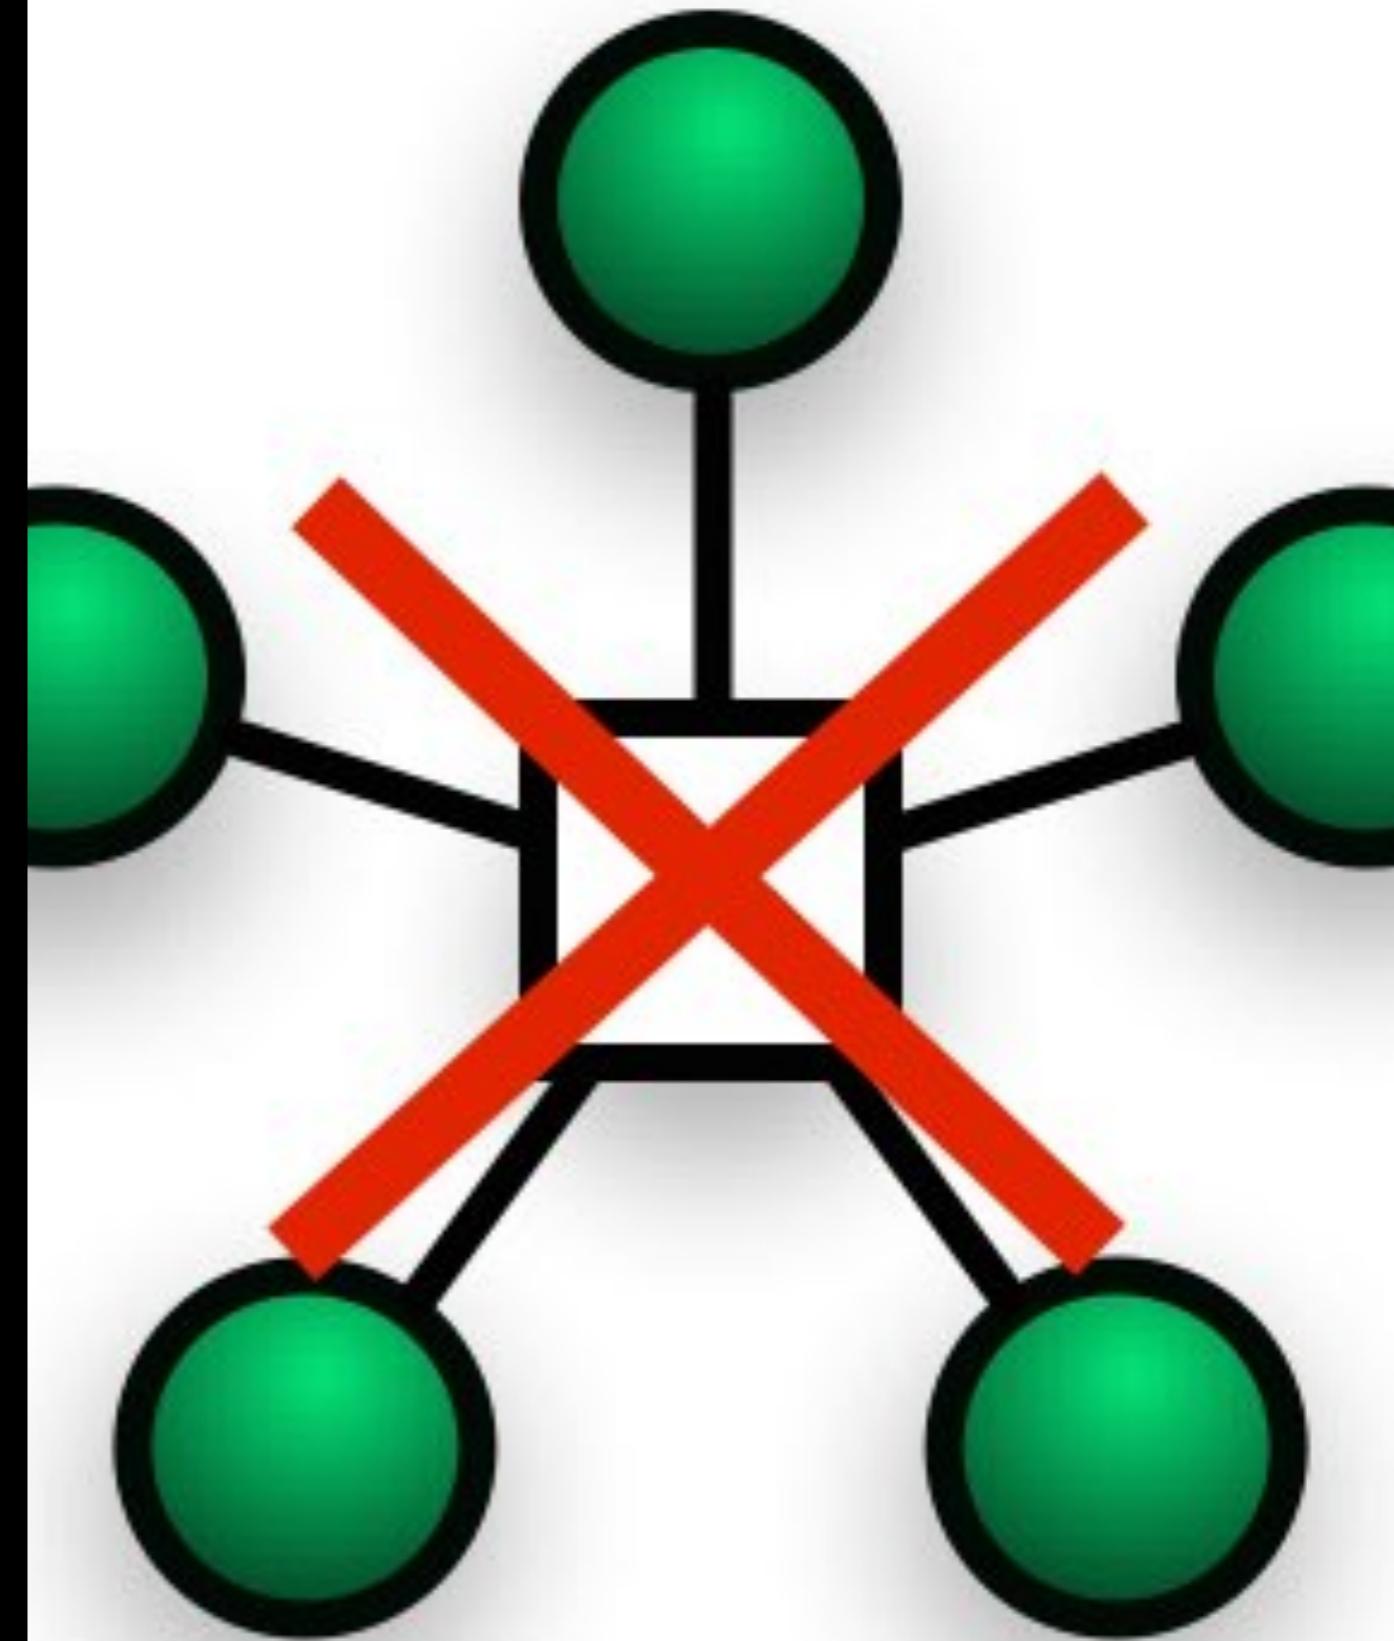
NOTE: this is a *fiction!*

²George Orwell, "Nineteen Eighty-Four", 1949.

Why cloud computing has become so *dystopian*? -- because we have sold freedom for convenience

- Ubiquitous/global accessibility
- Concentrated data for easy analysis
- Easy control of the information flow
- No extra cost for sharing
- No need to think about where the information locates

The
inconvenient
truth of
centralized
systems: what
if the core/
cloud fails?



Inconvenience of centralized systems

- Ubiquity or *no accessibility*
- When the core fails, no alternative
- When the core loses data, *no backup*
- The system performance is restricted by the capability of the core
- Endpoint systems will lost *all capabilities*

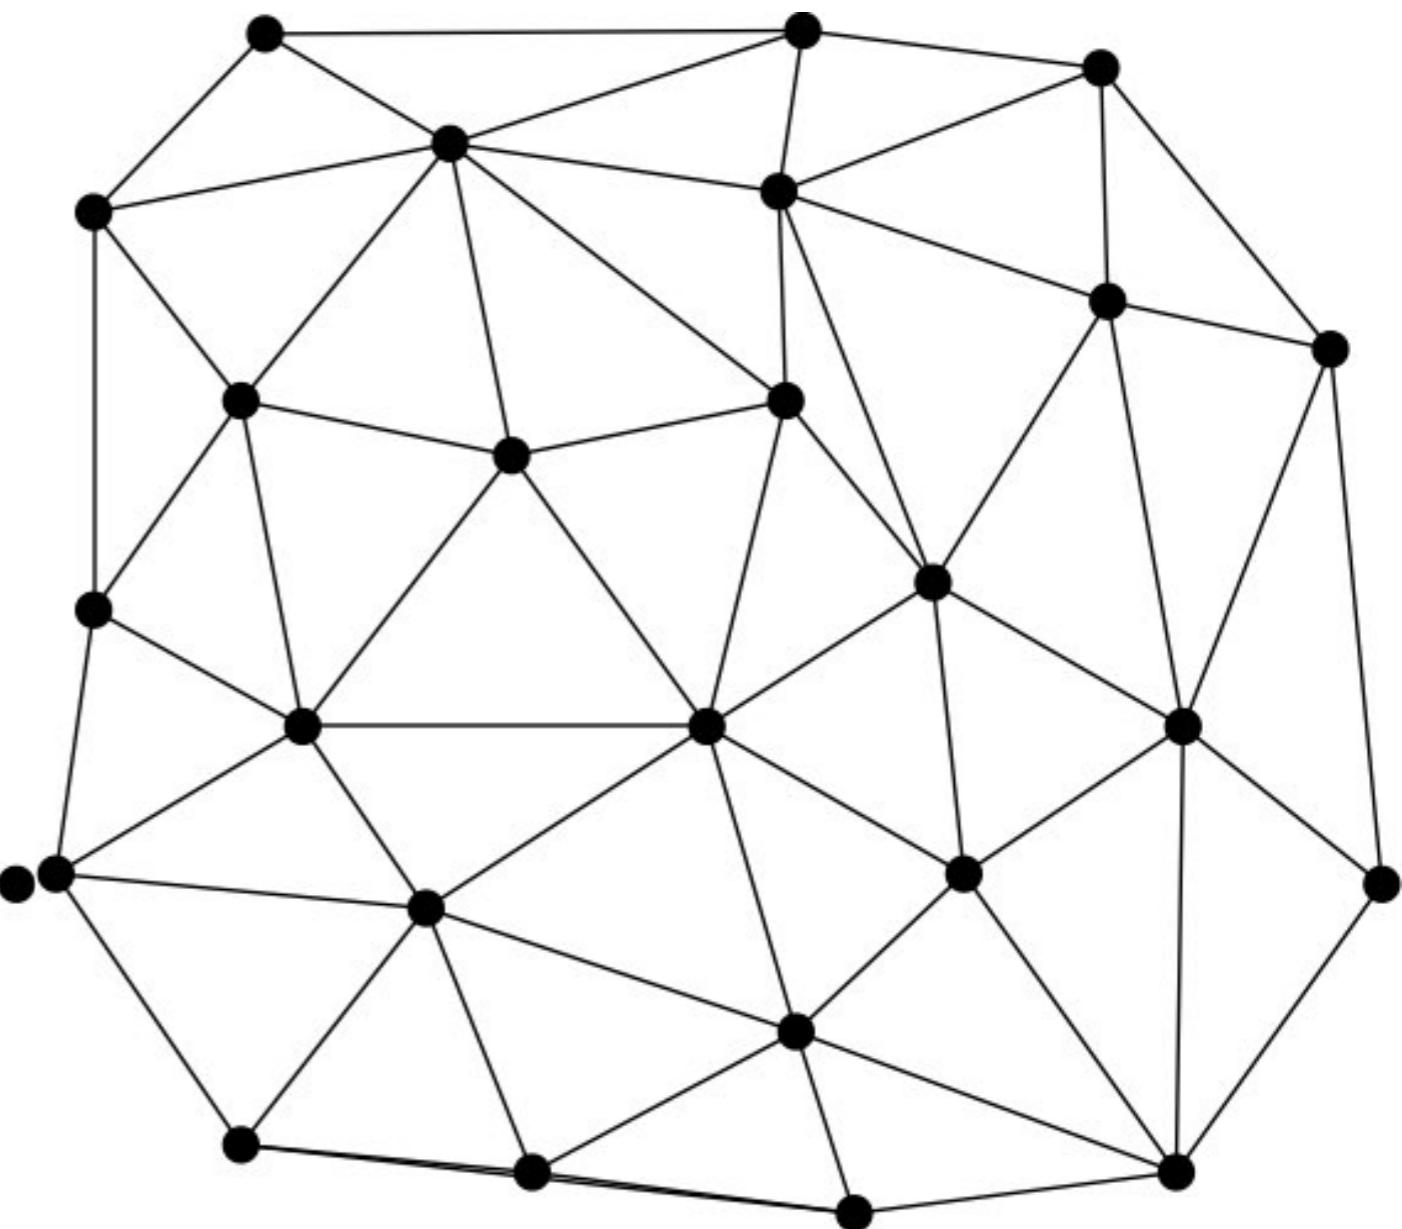
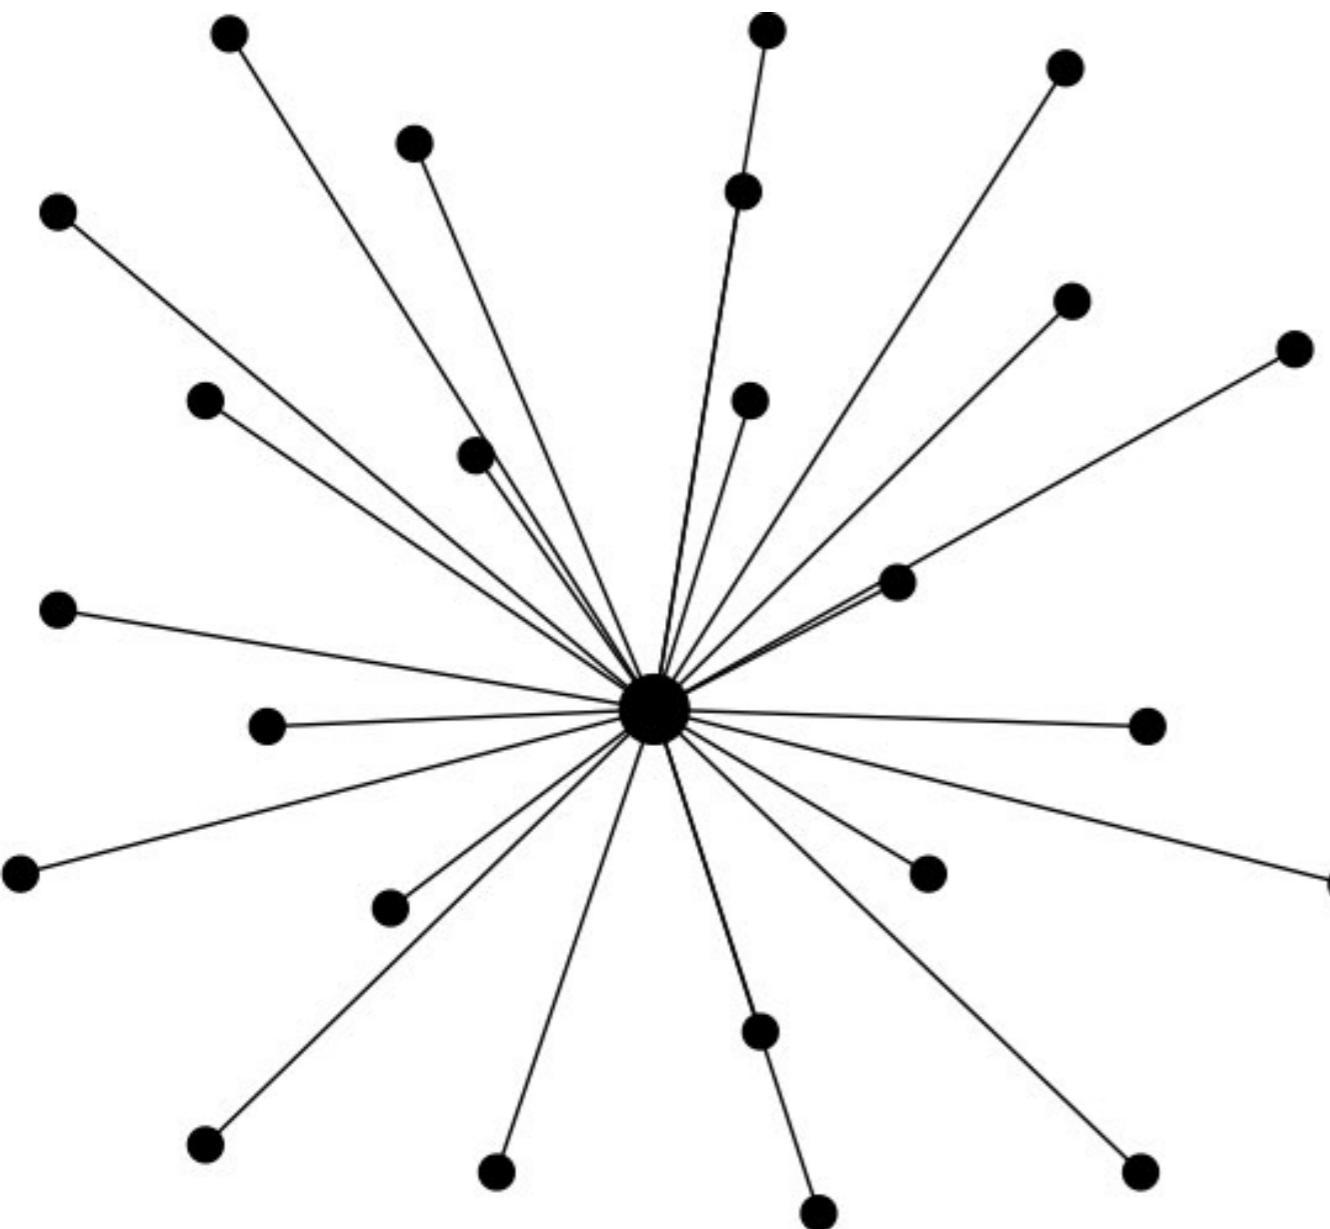
Centralized systems are *not* sustainable

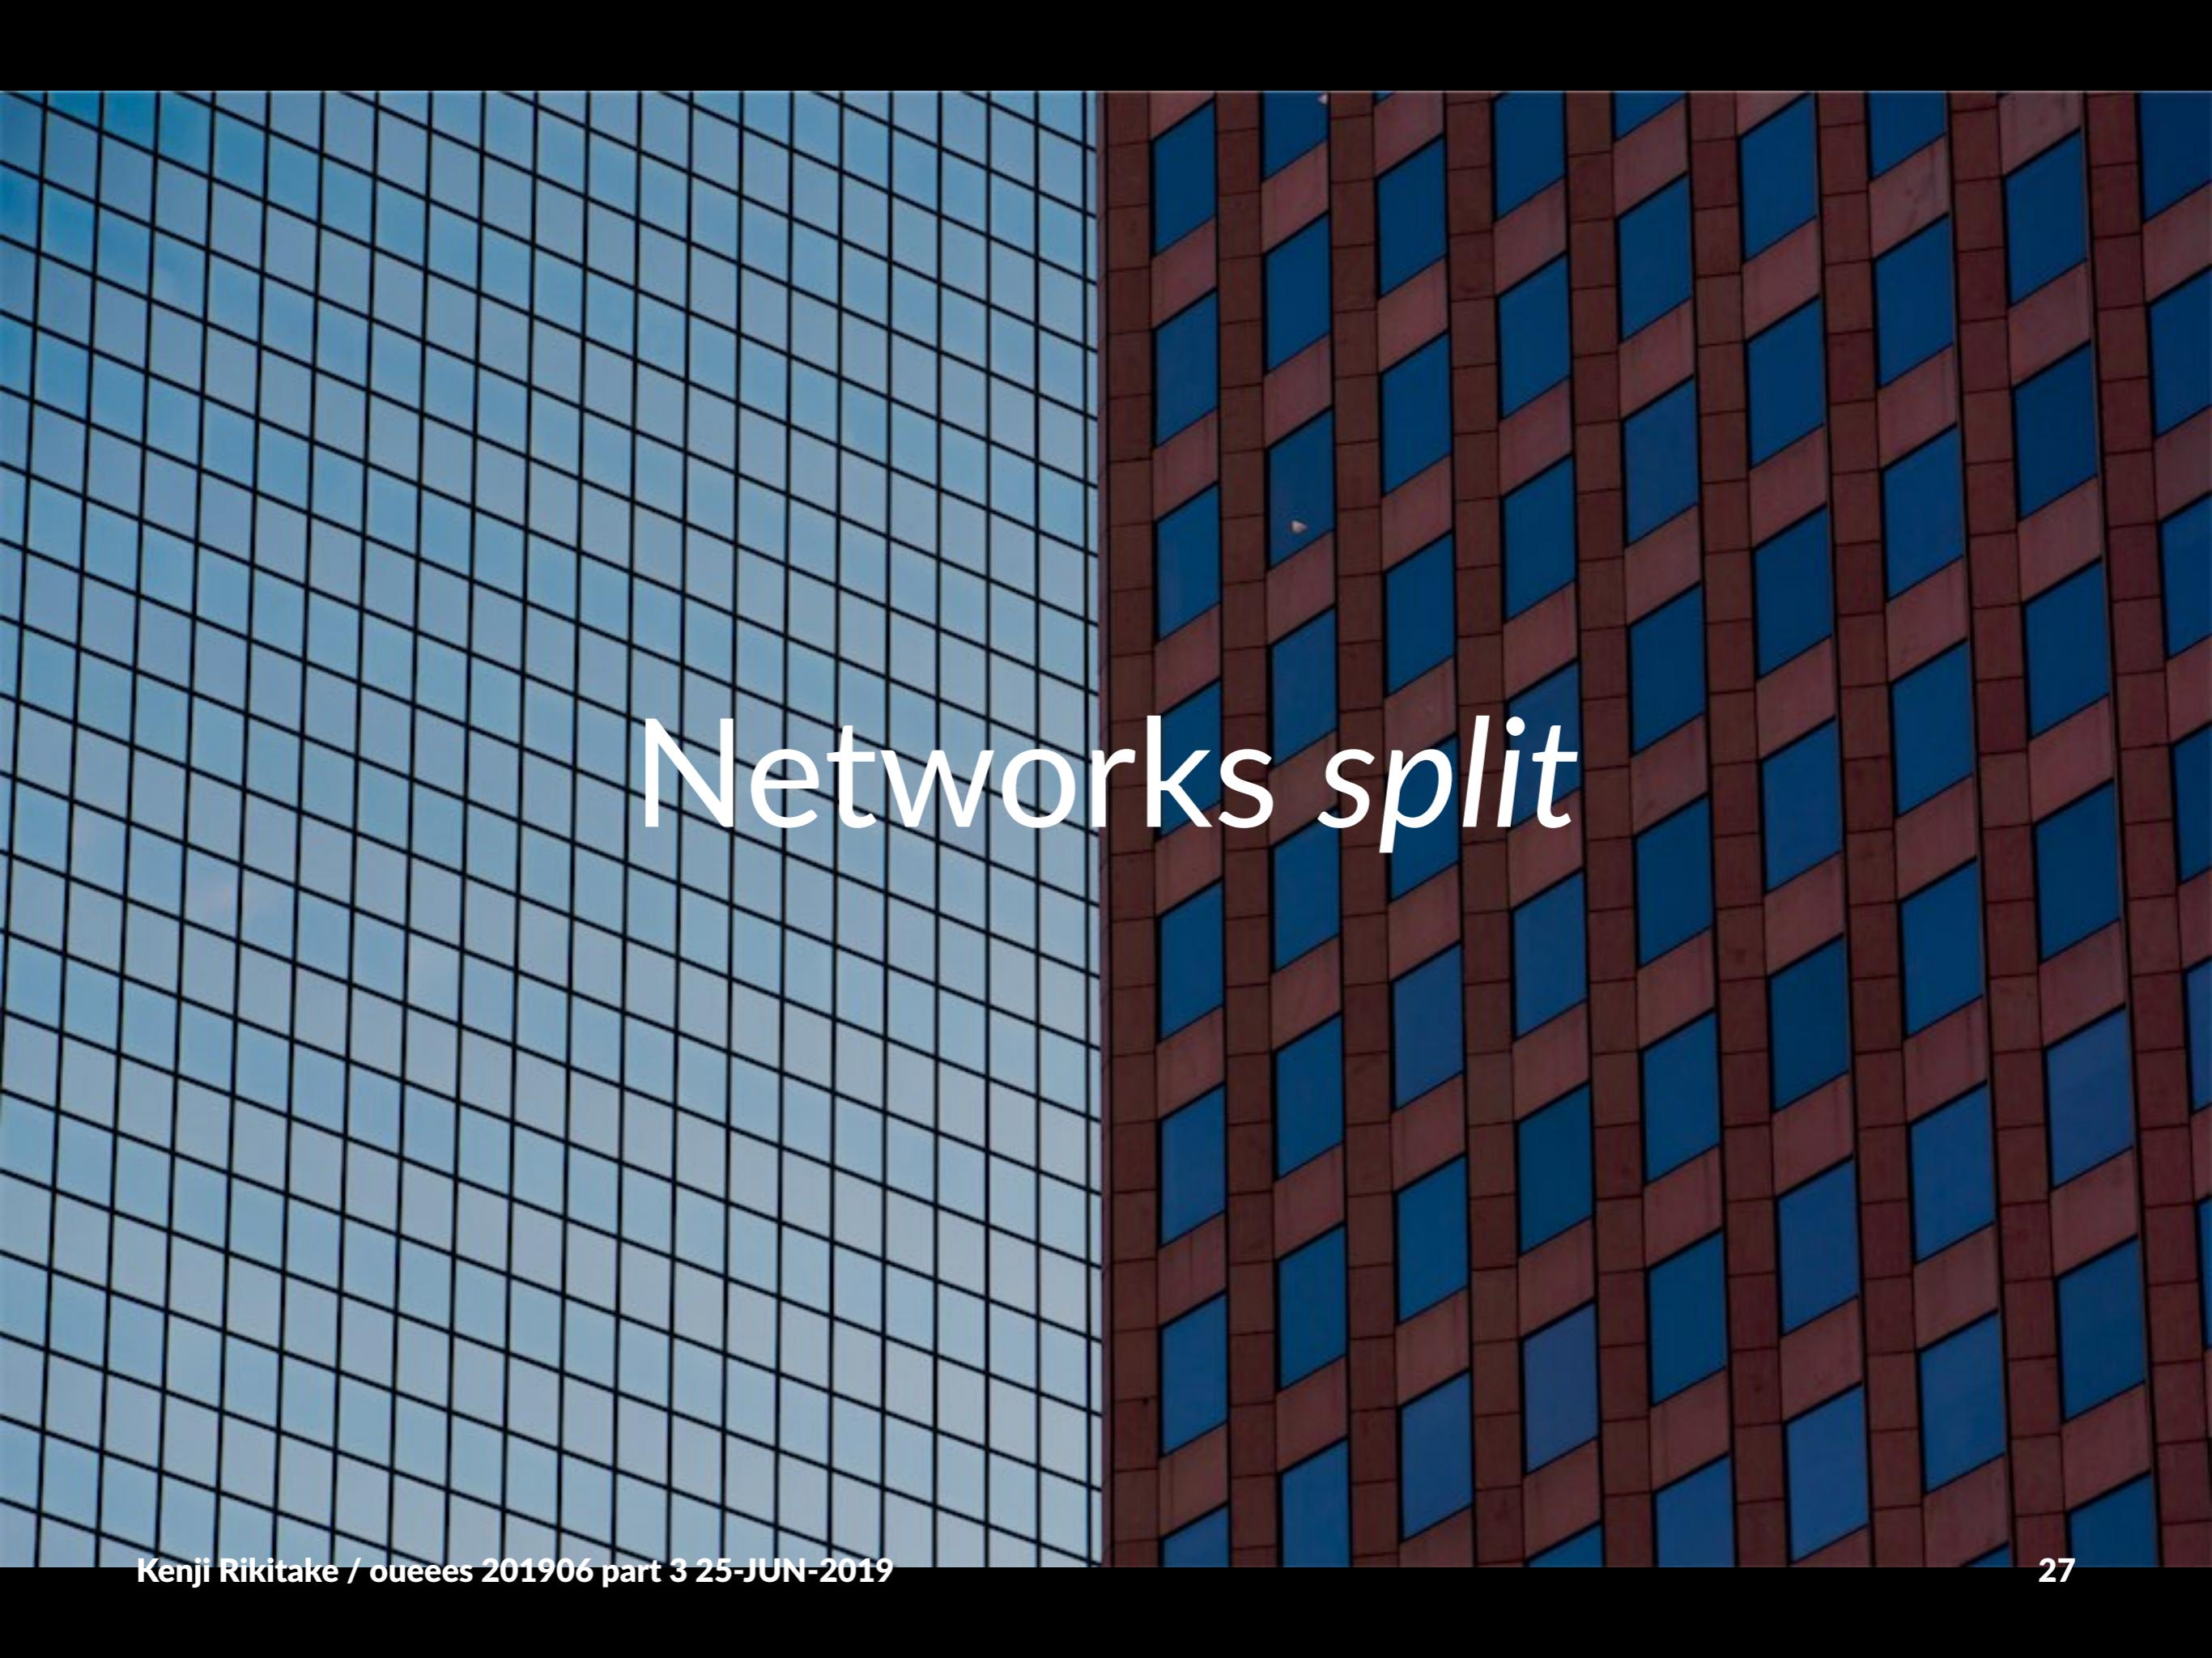
- ...then how cloud computing systems manages the sustainability?

Distributed systems provide
sustainability and resilience
against failures

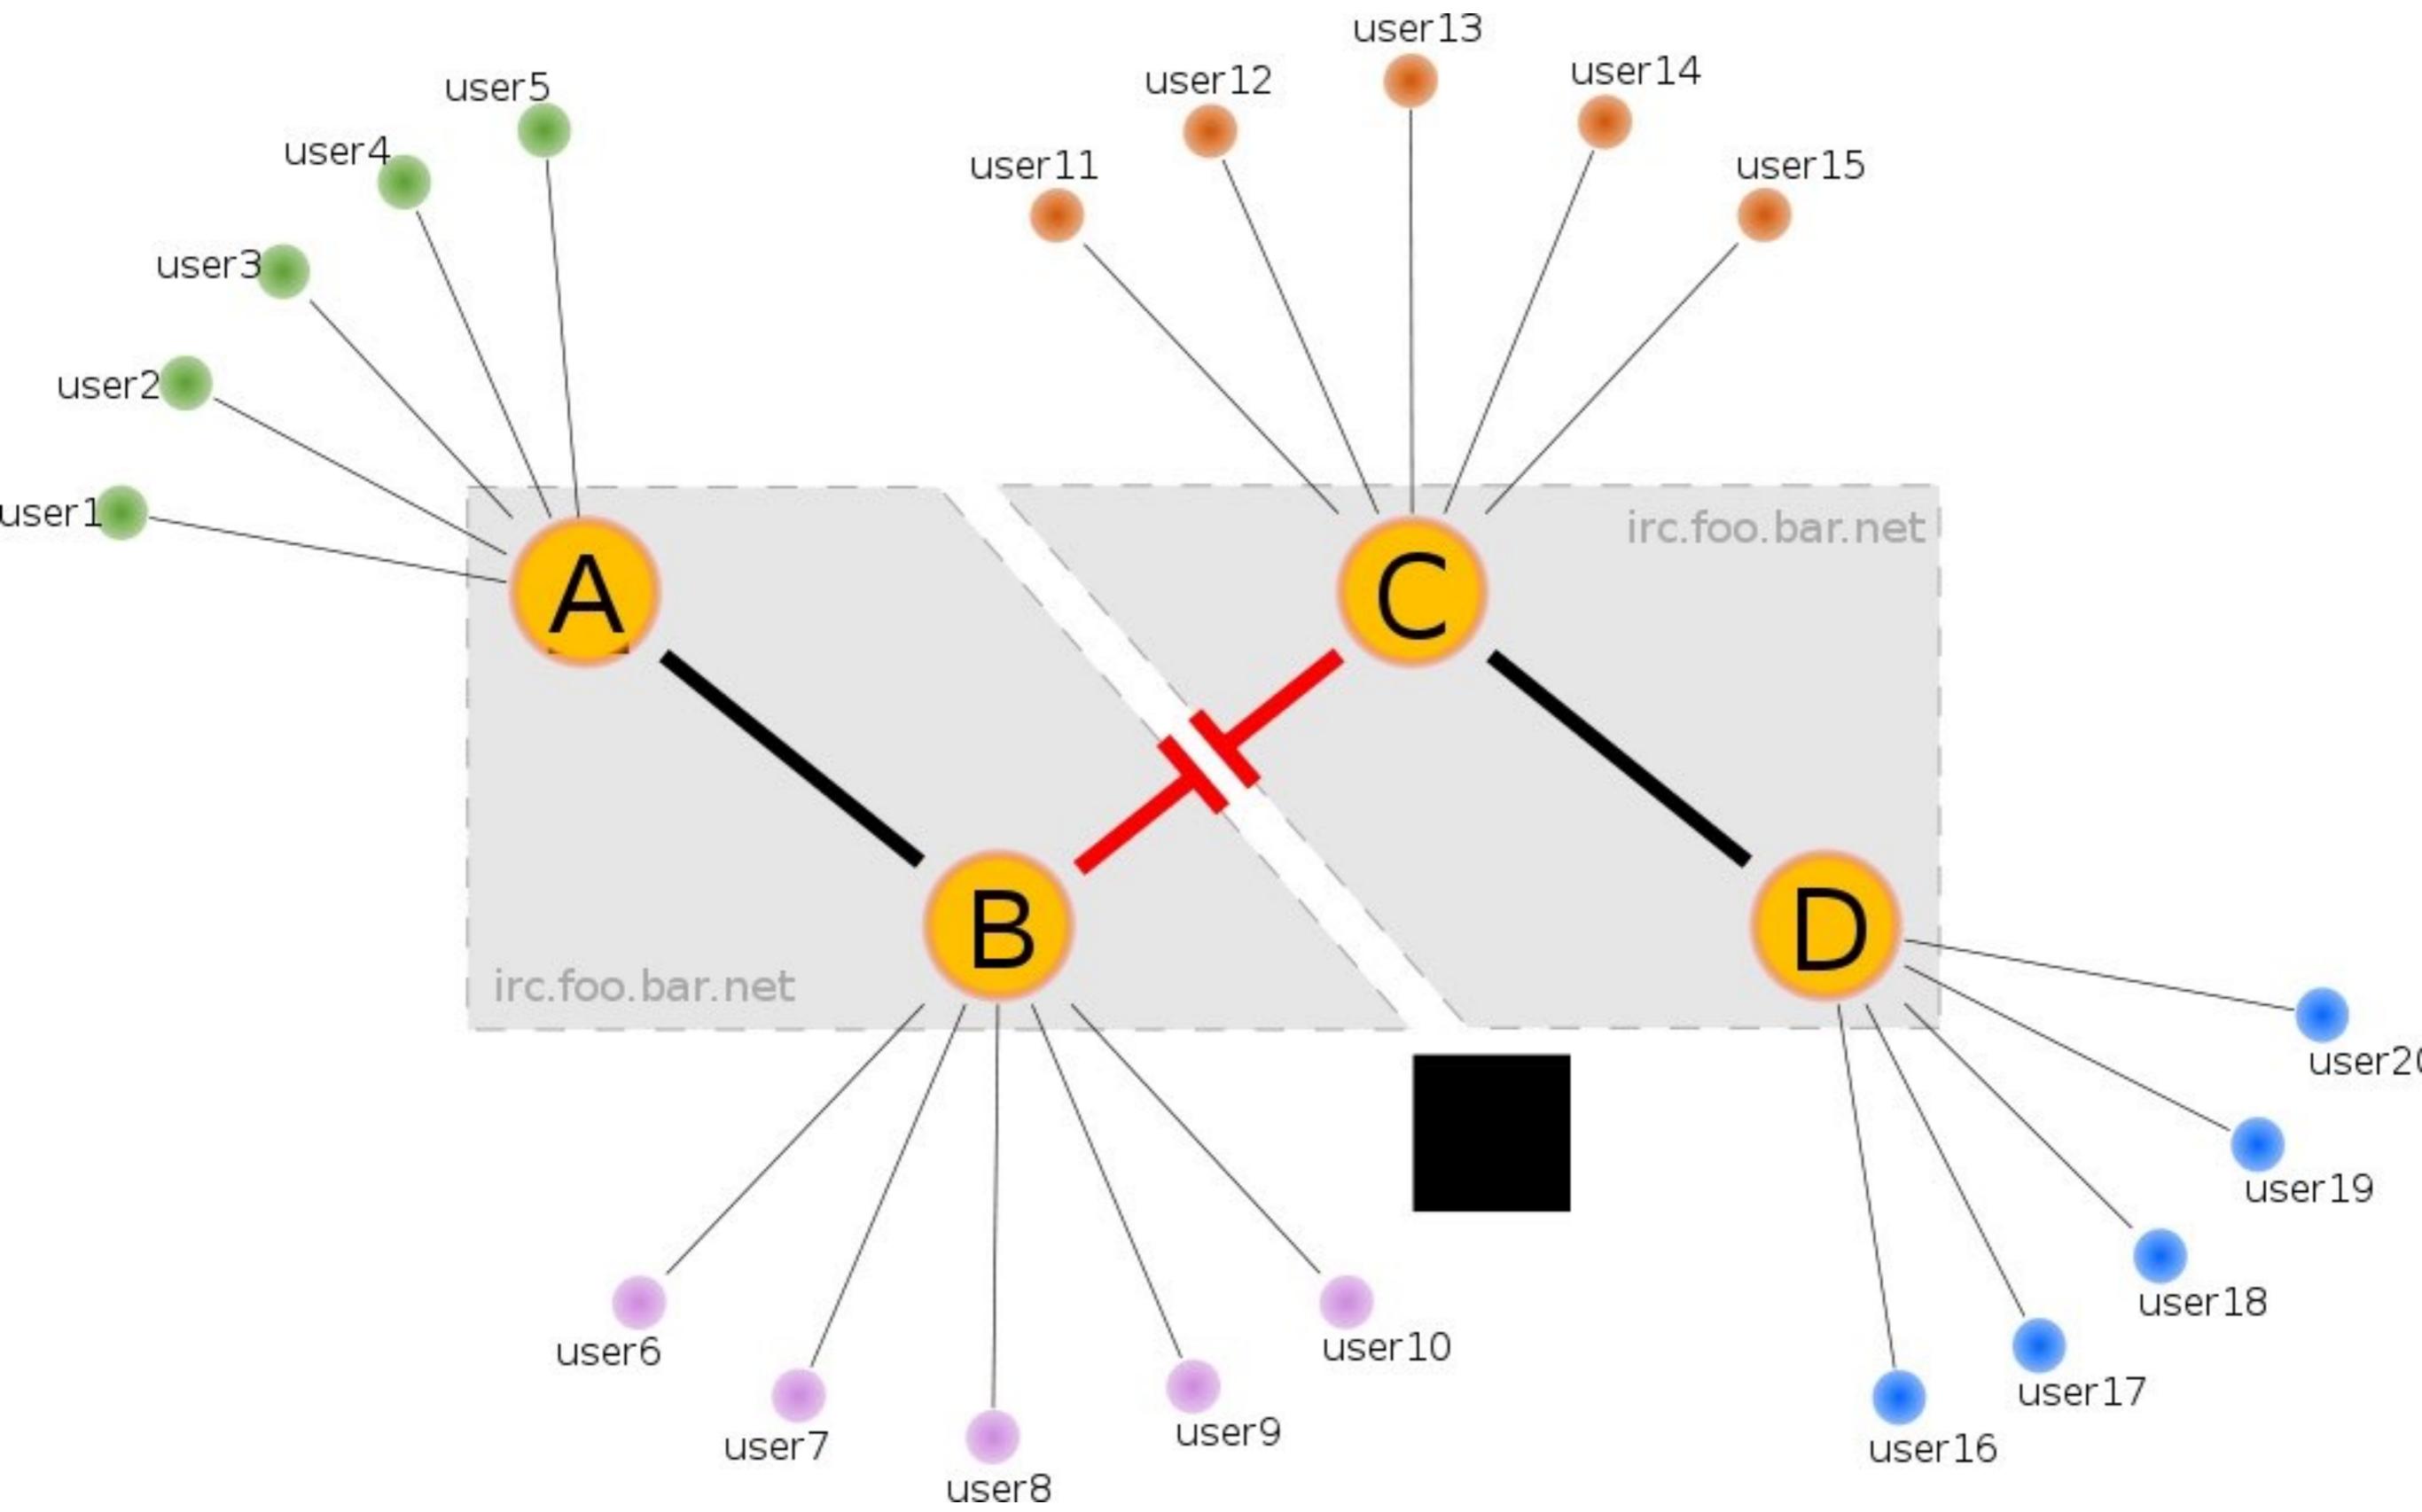
An aerial black and white photograph of a complex multi-level highway interchange. The image shows several intersecting roads with various ramps and overpasses. In the background, there are residential areas with houses and some industrial buildings. The word "Networks" is overlaid in large, white, sans-serif letters across the center of the image.

Networks

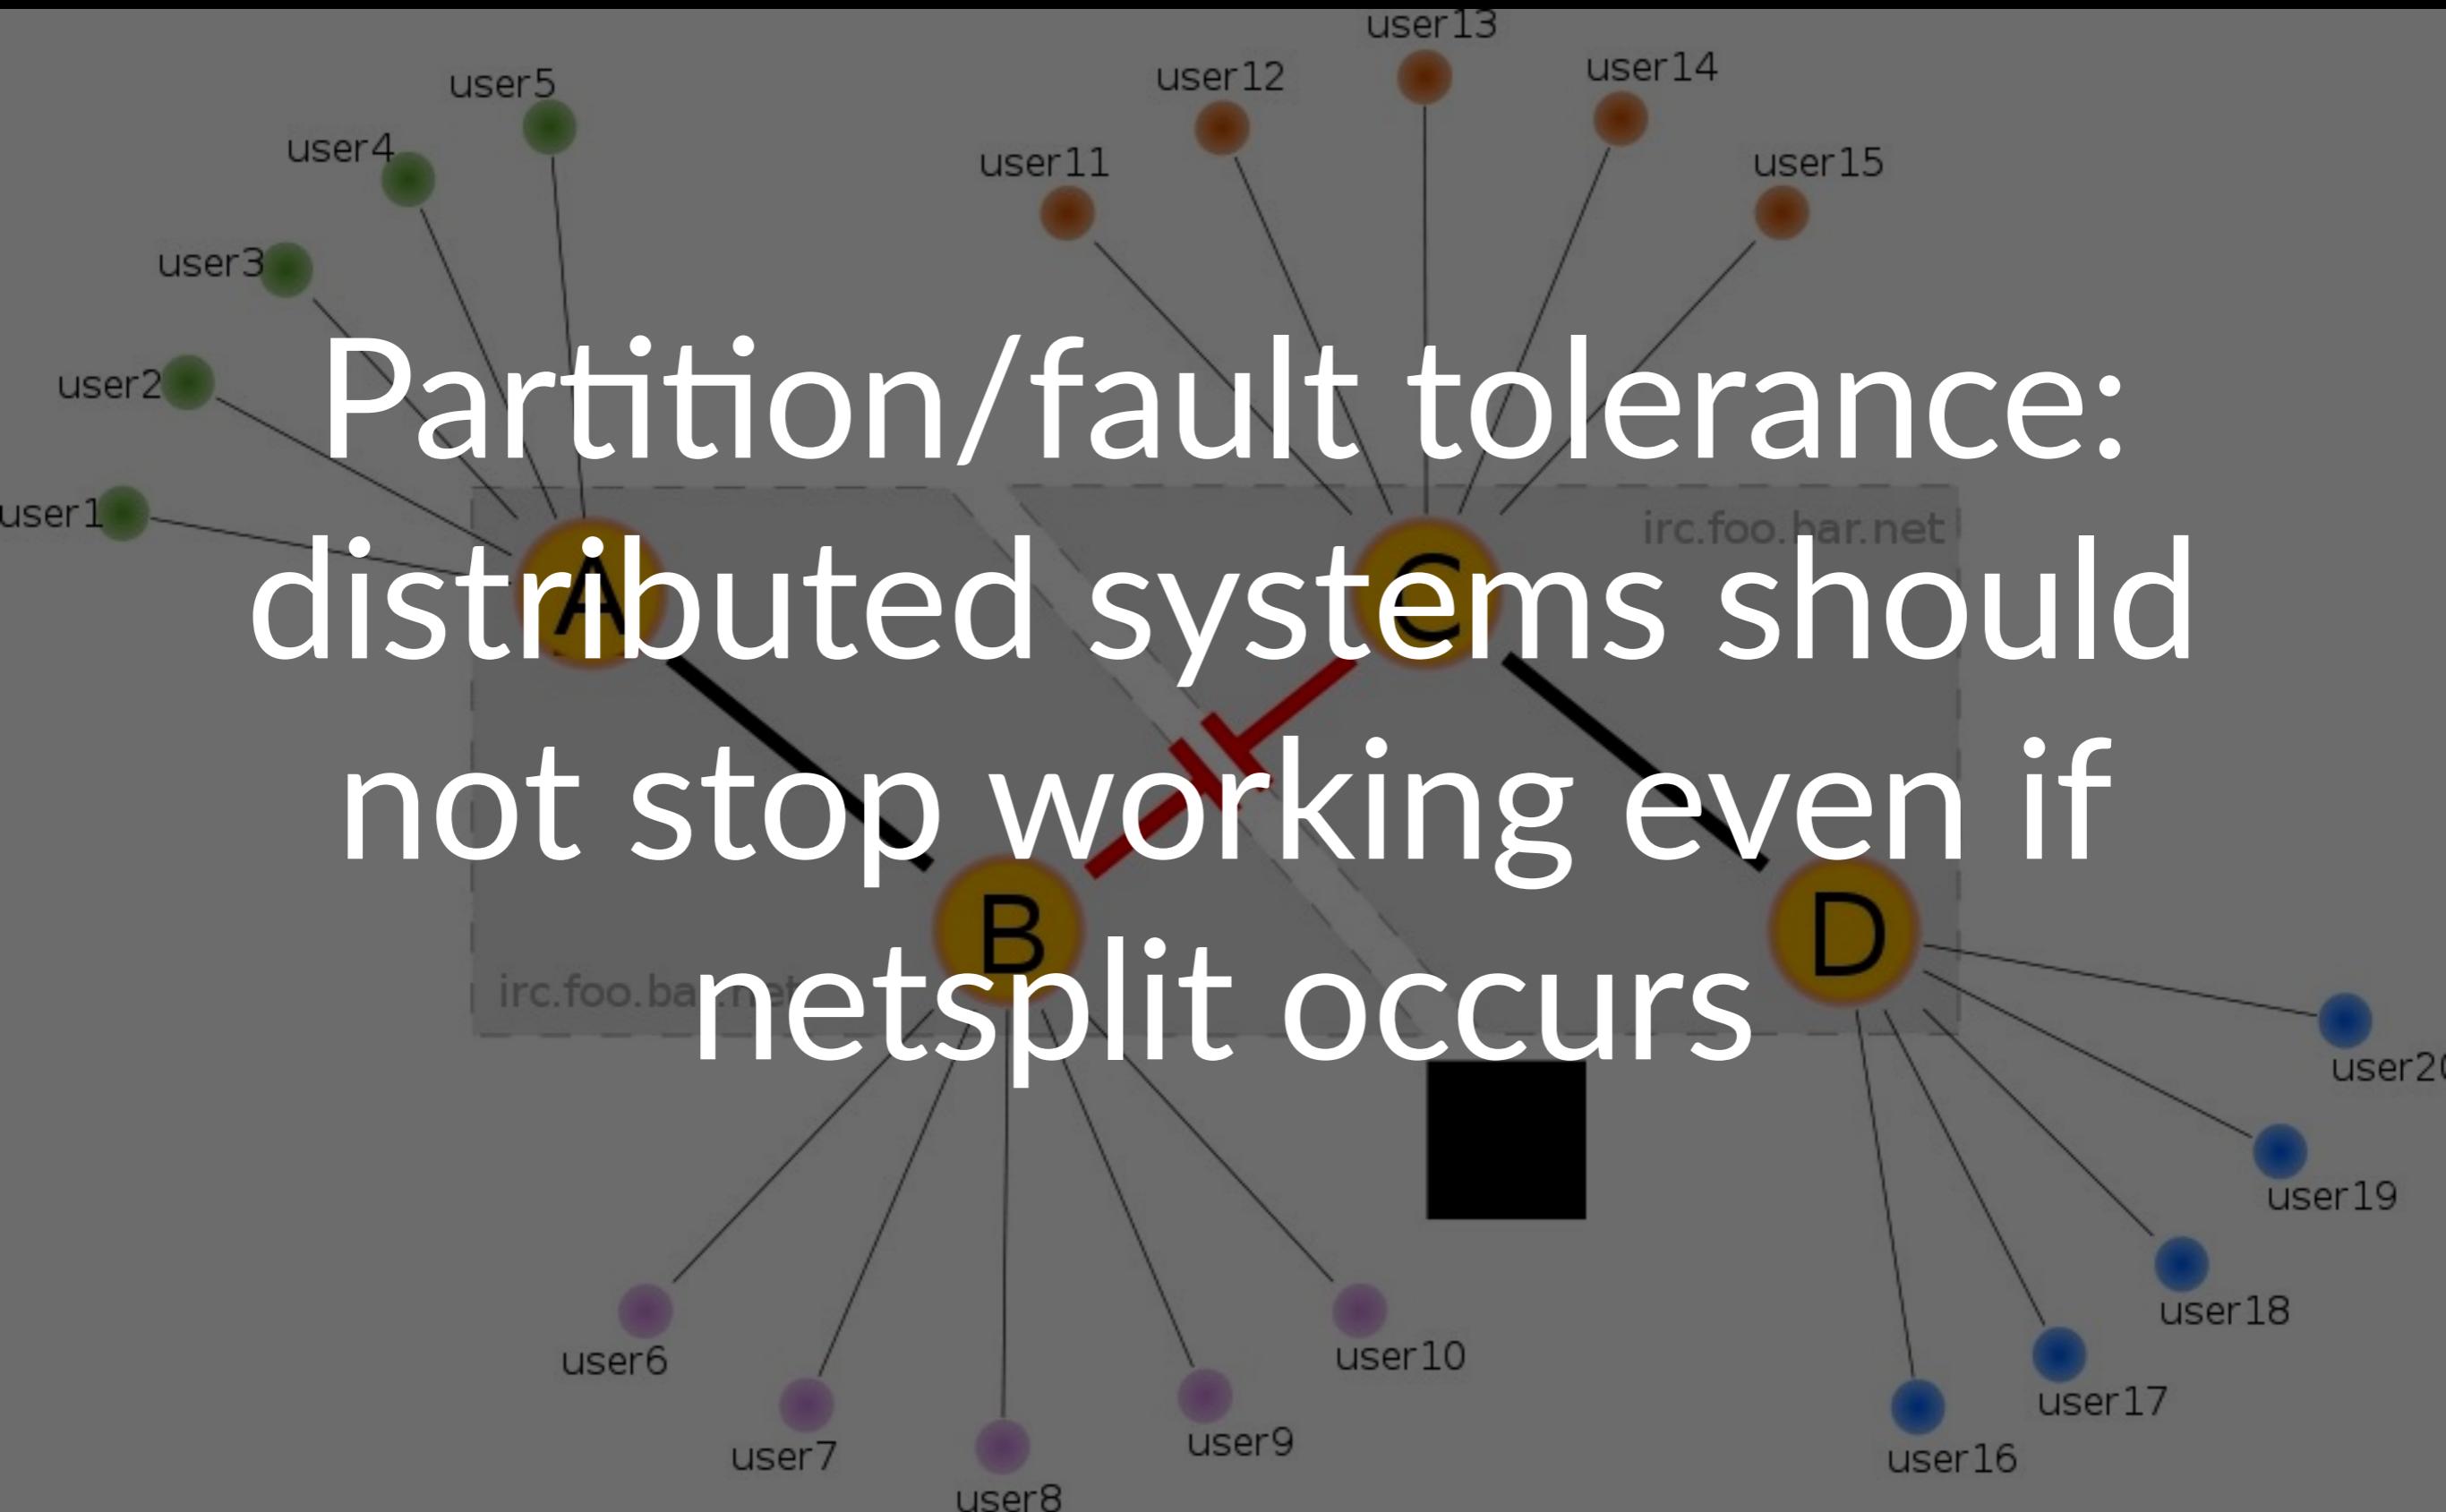




Networks split



Partition/fault tolerance:
distributed systems should
not stop working even if
netsplit occurs



Real-world challenges

- Natural disasters
- Device failures
- Human operation errors
- Political impediments
- Social resentments

Handling failures

- Redundancy: keeping backup units ready
- Fault tolerance: keeping systems running even the components fail
- Resilience by failing fast: early detection of failures and invocation of the recovery procedures

Why fault tolerance?

- Hard disk MTBF $\sim=$ 1 million hours
- 1000 hard disks running 24 hours \times 365 days = 8.76 million hours
- If you're running a system with 1000 hard disks, **9 out of 1000** will fail in a year
- Recovery of a disk content takes often *a day*; you can't stop a system for *a day*, can you?


Requirement to keep the systems fault tolerant

- Redundancy: two or more resources for each unit of processing
- Supervising the failure of the units by an independent supervisor
- Rollback capability: undo the incomplete operations and retry


Consistency issues of distributed systems

- Locking/synchronization: waiting all data to be ready to compute or proceed to next step
- Choosing the *right* data: which data is *correct*?
- Supervision: fault detection and restarting

Each function unit runs on
its own speed



Supervision
tree example



Try to
restart

when
crashed

Eight Fallacies of Distributed Computing³ (1/2)

- **The network is reliable**
- **Latency is zero**
- **Bandwidth is infinite**
- The network is secure

³ <https://blog.fogcreek.com/eight-fallacies-of-distributed-computing-tech-talk/>

Eight Fallacies of Distributed Computing (2/2)

- Topology doesn't change
- There is one administrator
- Transport cost is zero
- The network is homogeneous

Summary: centralized
computing is fragile;
distributed computing is
fault tolerant but hard

Appendix 1: references for further study


基礎からわかる TCP/IPネットワーク コンピューティング 入門 第3版

- ・ オーム社 ISBN
9784274050732
- ・ 達人出版会の電子書籍



プロフェッショナル IPv6

- ラムダノートの紙の本+電子書籍
- Boothの無料版



Appendix 2: on choosing your career and professionality

When I chose my career and professionalism?

- Age 9: computers and English
- Age 10 ham radio and electronics
- Age 14: writing commercial software
- *Age 23: finally decided to make my living on my computer software professionalism, with my English proficiency*

If I were at age 22, what I would do
after getting a Bachelor's degree?

- Get out of Japan ASAP
- Explore the computer skills
- Do something unpopular

Go abroad

What are the most important things to pursue engineering/scientist career?

- Physical strength
- Mental strength
- Curiosity

Curiosity matters

Photo credits

- All photos are modified and edited by Kenji Rikitake
- Photos are from Unsplash.com unless otherwise noted
- Title: NASA
- Modern Computing is Cloud Computing: Rayi Christian Wicaksono
- Cloud Computing: https://commons.wikimedia.org/wiki/File:Cloud_applications_SVG.svg, licensed under Creative Commons CC0 1.0 Universal Public Domain Dedication
- Intertwined network of computers: https://en.wikipedia.org/wiki/File:Cloud_Computing.jpg, licensed under Creative Commons CC0 1.0 Universal Public Domain Dedication
- Web services are clusters of computers: Kenji Rikitake, at Kyoto University ACCMS, April 2017
- Presidio Modelo Prison: By Friman [Public domain], from [Wikimedia Commons](#)
- Panopticon: Jeremy Bentham, Public domain, via [Wikimedia Commons](#)
- Contrasted Residences for The Poor: By Augustus Pugin [Public domain], via [Wikimedia Commons](#)
- Networks: Irina Blok
- Networks Split: Pietro De Grandi
- Netsplit: https://commons.wikimedia.org/wiki/File:Netsplit_split.svg, in public domain